
From Pull Requests to Sandboxed Restoration: A Logbook of Building an LLM Hardware-Agent Benchmark

David Andrews*, Saketh Reddy*
Georgia Institute of Technology
{dandrews47, sreddy347}@gatech.edu

Abstract

This logbook documents two and a half months of work building an LLM-agent benchmark for hardware description languages, ending with the RTL-UNIVERSE benchmark we report in our companion paper. The arc covers four prior repositories. The first was a GitHub pull request data pipeline (SWE-Universe-Verilog). The second was an early evaluation harness (RTLBench). The third was a port of NVIDIA’s CVDP problem suite into our infrastructure (CVDP-Harbor). The fourth was a spec-to-GDSII benchmark we built and iterated weekly through three pivots (Spec2GDSBench). The current state, RTL-UNIVERSE, takes a different approach by framing tasks as project-scale restoration of open-source HDL projects. Although the project began as a CS 3220 research option, we treat the final submission as a benchmark-design study for LLM hardware agents, with RISC-V cores, AI accelerators, and real open-source RTL projects serving as the evaluation substrate. The logbook records what we tried, what failed, what we learned, and why we pivoted, with concrete references to commits, weekly reports, and run results in the project’s `other_repos/` directory.

1 Goal

The high-level question we set out on was the following. Can current LLM agents do real hardware-engineering work, and how do we measure it? The literature on hardware LLMs in early 2026 was almost entirely module-level Verilog generation [1–4], with leading models reporting 90 percent or higher pass rates on benchmarks that ask for “a 4-bit counter from a description”. That left two open questions. How would those agents fare on the rest of a real chip-design workflow (build systems, integration, testbenches, physical design), and is there a benchmark that we trust to grade them on it?

Over two and a half months we built five repositories, in roughly chronological order: SWE-UNIVERSE-VERILOG, RTLBENCH, CVDP-HARBOR, SPEC2GDSBENCH, and RTL-UNIVERSE. Each one taught us something that we then either folded into the next or used as the reason to start from scratch. Table 1 lists each phase, the key commits, the result, and the decision we made next. The remainder of this document expands each row.

2 March 2 to 6: SWE-Universe-Verilog (data pipeline)

What we tried. SWE-bench [5] grades agents on real GitHub bugs in real Python projects. The first thing we tried was the same recipe but for HDL. The SWE-UNIVERSE-VERILOG pipeline crawls GitHub for repositories whose primary language is Verilog or SystemVerilog, or that match

*Equal contribution. Authors listed in alphabetical order.

Table 1: Project timeline. Each row is a phase, the repository it lived in, the commit range that bounds it, the headline result, and the decision that ended it.

Period	Repo	Commits	Result	Decision
Mar 2–6	swe-universe-verilog	47f1b40 .. 1b6b034	401 strict-filter PRs from 3,296 HDL repos.	PR-replay does not isolate cleanly. Pivot to CI-graded restoration.
Mar 2–3	rtlbench	45e61b9 .. 38a68a1	First Claude runs. Toolchain-version mismatches and ambiguous testbenches.	Adopt per-task Dockerfiles. Require mechanically-checkable verifiers.
Mar 7–13	cvdp_benchmark	c002eb5 .. eba655b	3 Claude variants, 90+% pass on CVDP.	Module-level HDL is saturated. Move to project scale.
Mar 23–28	spec2gdsbench (W1)	9821882 .. d96e177	Sonnet 0.964 mean over 64 runs on 8 tasks.	Near-saturated. Pivot to a single SoC task.
Mar 31–Apr 6	spec2gdsbench (W2)	ba338b7 .. 2fa0d8d	Sonnet best 0.616 on RV64GC Linux-boot SoC.	Single task gives poor signal. Pivot to a graduated ladder.
Apr 7–12	spec2gdsbench (W3)	723f8c3	3 AI-accelerator tasks pass full GDS flow as references.	Specification authoring is the bottleneck. Pivot to upstream restoration.
Apr 23–May 8	rtl-universe	1024c2a .. current	7-model $\times$ 11-task matrix; audit reveals cheating.	Companion paper. Continue.

HDL topics (verilog, fpga, rtl, hdl, asic). It collects all merged pull requests, then applies a multi-stage filter inspired by the SWE-Universe data-curation methodology [6].

Pipeline yield.

Stage	Count
Repos discovered (Verilog/SV on GitHub)	3,296
PRs collected (merged)	173,106
After heuristic filter (size, bot, merge)	27,446
After testbench detection	5,282
After patch separation (test + fix files)	3,070
Quality filter (relaxed)	1,241
Quality filter (strict, with issue linkage)	401

What we learned. Even after aggressive filtering, the yield from the GitHub HDL ecosystem is roughly an order of magnitude lower than the Python equivalent. 401 strict-filter PRs across 3,296 repositories is enough to start a benchmark, but each one had to be curated by hand to verify that it isolated a single change with a real before-and-after test. We built RTL BENCH (Section 3) on top of this data and quickly discovered that PR-based hardware bugs do not in practice isolate cleanly. Most are part of a larger refactor, depend on FPGA-vendor toolchains, or test only a manually-asserted simulation trace.

This phase informed two decisions in RTL-UNIVERSE. We chose to abandon PR-replay as the framing, and we chose restoration (delete-and-rebuild) instead. Restoration gives the agent a real test suite (the upstream’s own CI) without requiring us to curate a clean before-and-after diff.

3 March 2 to 3: RTL BENCH (first agent runs)

What we tried. Using the SWE-UNIVERSE-VERILOG filtered set, plus a few additional integration projects (corundum networking, CVA6 RISC-V core), we built a quick evaluation harness called RTL BENCH. It clones the upstream repo at the parent commit of a chosen PR, sets up the toolchain in a container, and asks Claude to make the change. The early commits were exploratory (“init”, “snap”, “first claude run”, “claude plan expand not run yet”, “ran full eval”, commits 45e61b9 through 38a68a1).

What we learned. The first runs worked just well enough to show two failure modes that ended up shaping everything that followed. The first was toolchain fragility. Verilog projects often pin specific Verilator and Yosys versions, and the agent would happily install a different one and write code that compiled in the wrong version. We ended up writing per-task Dockerfiles, which is what the next two projects (CVDP-Harbor and Spec2GDSBench) inherited. The second was test-suite ambiguity. Many of the PRs in our filtered set have testbenches that “run” on master but actually only pass when

accompanied by a manually-checked waveform. The agent does not know this. We needed verifiers whose pass criterion is mechanically checkable. This is part of why RTL-UNIVERSE restricts itself to projects with binary-pass CI.

4 March 7 to 13: CVDP-Harbor

What we tried. At this point NVIDIA released the Comprehensive Verilog Design Problems (CVDP) benchmark [7], a 783-problem suite of mostly module-level Verilog tasks with both non-agent and agent variants. Rather than re-invent the problems, we ported CVDP into *Harbor*, our internal evaluation harness used by the rest of the projects. The conversion tool is at `cvd_benchmark/` commit `c002eb5`, with three more rounds of fixes and reruns over the next week.

Result. We ran three Claude variants on the converted suite. Pass rates clustered above 90 percent (`cvd_benchmark` commit `891039d`), consistent with what NVIDIA themselves reported. This was when we became confident that module-level Verilog generation is saturated for current frontier models, and that the open question is at the project scale. At this point we stopped working on CVDP and started a benchmark we hoped would not saturate.

5 March 23 to April 12: Spec2GDSBench (three weekly pivots)

Goal. Given a behavioural specification and a test suite inside a containerised environment with the full open-source ASIC toolchain (LibreLane and the SkyWater Sky130A PDK), ask the agent to autonomously produce a DRC-clean, LVS-verified GDSII file. The grading funnel is functional simulation passes, then Yosys synthesis, then LibreLane place-and-route, then KLayout and Magic DRC, then Netgen LVS, then GDSII export. Each stage gates the next.

Week 1 (March 24 to 28): eight isolated tasks. The first version had eight tasks: `gcd_unit`, `uart_tx`, `aes128_encrypt`, `femtorv32_rv32i`, and four AI accelerator primitives (`conv1d_engine`, `maxpool2d`, `requantize_pipeline`, `matmul_4x4_int8`). 64 runs of Claude Code Sonnet at \$5 per run produced mean reward 0.964 in normal mode and 0.829 in blind mode (no test suite visible). All 32 normal-mode runs passed all functional tests. The only score loss was timing closure during physical design [8]. The benchmark was near-saturated, and we pivoted.

Week 2 (March 31 to April 6): single SoC, Linux boot. We replaced the eight tasks with one task. Design an RV64GC RISC-V SoC capable of booting Linux. The agent receives only the external interface (AXI4 and UART) and a software contract (108 ISA compliance tests, 10 bare-metal programs, OpenSBI and BusyBox boot via QEMU’s virt machine convention). Claude Sonnet 4.6 across five runs scored a best of 0.616 (85 of 108 ISA, 10 of 10 bare-metal, OpenSBI boot). A 15-hour run reached 99 of 108 ISA and showed partial Linux boot during debugging [9]. This was a much harder benchmark, but a single task gave poor signal because an agent either built a working SoC or scored near zero. We pivoted again.

Week 3 (April 7 to 12): AI accelerator integration ladder. We then designed three new tasks at graduated difficulty, all with AXI4 interfaces: `systolic_array` (8×8 INT8 MAC array, AXI4-Lite MMIO and AXI4-Stream), `conv2d_engine` (3×3 convolution with line buffer, AXI4 master), and `dma_engine` (memory-to-memory DMA, AXI4-Lite and AXI4 master). All three reference solutions pass the full ASIC flow through GDSII. Only the DMA meets timing at 50 MHz [10]. We were partway through getting agent runs on these when the next pivot happened.

What we learned across the three weeks. Specifications are hard. Every week spent on Spec2GDSBench, we spent some fraction of the time on whether the spec was over-constraining the agent (forcing a specific microarchitecture) or under-constraining it (letting the agent solve a different problem). Real industrial specs are often closer to “here is a JIRA epic with 30 linked tickets” than “here are the AXI4 signal names”. We could not write a clean spec for a realistic design.

Build systems matter as much as RTL. The agent’s score on Spec2GDSBench was dominated by physical design rather than RTL design. Timing closure, place-and-route choices, and LibreLane configuration all dominated the score. We wanted RTL-design ability to be the bottleneck.

Real projects already have all of this. The thing that finally moved us was realising that lowRISC’s Ibex, NVIDIA’s NVDLA, OpenPiton, and the rest already have testbenches, build systems, CI configurations, and reference implementations carefully tuned by their maintainers. We do not need to write a spec for a UART because lowRISC has a real one and grades it with their own `simple_system`. We just need to delete files.

6 April 23 to May 8: RTL-Universe

What it is. RTL-UNIVERSE is the result of all the prior decisions. Real upstream projects, real upstream CI, no spec to write, restoration as the framing. The task list is in our companion paper. The harness builds on the Harbor runner that started in CVDP-Harbor (Section 4). The Docker patterns follow Spec2GDSBench’s three-layer image. The verifier conventions are upstream’s CI.

The first three weeks. Most of late April and early May was task engineering. We wrote a `harbor-task-gen` skill (commits 1ad3b4d, 82ce12f, 1df25f5) that converts an arbitrary upstream HDL repo into a Harbor task with correct file-strip manifest, environment Dockerfile, and verifier script. The skill went through five iterations to handle five different upstream build systems (FuseSoC, Bazel, Make+Verilator, Icarus, cocotb). We added Ibex, Caliptra, PULP common_cells, secworks-aes, VeeR EL2, CVA6, Coral NPU, NVDLA, and OpenPiton. Roughly half of the tasks we tried did not make the final eleven. Some had upstream build systems that the agent could not drive at all (Caliptra). Some had test suites with too much state for the verifier to be deterministic (CVA6 in some configurations).

The runs. The headline run on the eleven tasks took roughly four days of wall-clock and seven model and harness combinations: Opus 4.7 and Sonnet 4.6 via Claude Code, GPT-5.5 (effort `xhigh`) via Codex CLI, and DeepSeek V4-Pro, GLM-5.1, Kimi K2.6, Qwen 3.6-27B via OpenCode and OpenRouter. The harness needed several mechanisms we had not anticipated when we started: per-account concurrency caps because Claude Code rate-limits per subscription, a disk supervisor because a Sonnet run leaked 812 GB of Verilator build cache through a bind-mounted credentials directory and would have crashed the lab server, and a stuck-job killer because the agents sometimes write at human pace for hours and then freeze.

The result, briefly. The exploratory baseline records 6 perfect scores for GPT-5.5 and 5 for Opus 4.7. The smaller open-weight models cluster at 0 to 2 perfect scores each. Reading the transcripts we found that 10 of the 18 perfect scores came either from sandbox bypasses (CDN mirrors, IP-literal DNS bypass, vendor-side fetch tools) or from two upstream verifier flaws (an exit-code-only test in Ibex and an `assert True` scoreboard in VeeR-EL2). After zeroing those cells, the strongest agents pass 2 of the 8 trusted tasks. We frame the matrix as an exploratory baseline plus an audit finding rather than as a definitive leaderboard. The full discussion is in the companion paper.

7 What we would do differently

Three lessons we would carry into the next iteration.

Sandbox first. We built the harness and added the DNS sinkhole later. By the time we audited the transcripts, most of the perfect scores had already been bypassed. A default-deny egress firewall with explicit allowlists, plus an audit pass that flags external network egress in real time, would have caught everything in this paper before the run completed.

Verifier hardening should be part of task acceptance. For every task we adopt from an upstream project, “does the agent score 1.000 with a stub” should be a check we run before greenlighting the task. Two of the three flagged tasks were exploitable in 17 minutes of agent time. We could have caught both during task setup.

Restoration is the right framing for now. Of the four prior approaches (PR-replay, CVDP module synthesis, Spec2GDSBench, restoration), restoration is the only one that simultaneously gives the agent a real test suite, avoids the specification problem entirely, scales to project-scale work without us having to write build infrastructure, and produces a benchmark whose top-line we trust. RTL-UNIVERSE as it stands is not the right benchmark forever, but the restoration framing has more headroom than what we had before.

Acknowledgements. We thank the lowRISC, chipsalliance, NVIDIA, Princeton, PULP, Google Coral, and secworks teams for the upstream projects this benchmark is built on. Repository contributions over this arc are split roughly evenly between the two authors, with several small contributions appearing under our other GitHub handles (Brojojo and spyre respectively, plus the SpiritSeal handle that was used during the early SWE-Universe-Verilog and RTL Bench work).

References

- [1] Mingjie Liu, Nathaniel Pinckney, Bruce Khailany, and Haoxing Ren. VerilogEval: Evaluating large language models for verilog code generation, 2023. URL <https://arxiv.org/abs/2309.07544>.
- [2] Yao Lu, Shang Liu, Qijun Zhang, and Zhiyao Xie. RtlM: An open-source benchmark for design rtl generation with large language model, 2023. URL <https://arxiv.org/abs/2308.05345>.
- [3] Shailja Thakur, Baleegh Ahmad, Zhenxing Fan, Hammond Pearce, Benjamin Tan, Ramesh Karri, Brendan Dolan-Gavitt, and Siddharth Garg. Benchmarking large language models for automated verilog rtl code generation, 2022. URL <https://arxiv.org/abs/2212.11140>.
- [4] Shang Liu, Wenji Fang, Yao Lu, Qijun Zhang, Hongce Zhang, and Zhiyao Xie. RtlCoder: Outperforming gpt-3.5 in design rtl generation with our open-source dataset and lightweight solution, 2025. URL <https://arxiv.org/abs/2312.08617>.
- [5] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. Swe-bench: Can language models resolve real-world github issues?, 2024. URL <https://arxiv.org/abs/2310.06770>.
- [6] Mouxiang Chen, Lei Zhang, Yunlong Feng, Xuwu Wang, Wenting Zhao, Ruisheng Cao, Jiaxi Yang, Jiawei Chen, Mingze Li, Zeyao Ma, Hao Ge, Zongmeng Zhang, Zeyu Cui, Dayiheng Liu, Jingren Zhou, Jianling Sun, Junyang Lin, and Binyuan Hui. Swe-universe: Scale real-world verifiable environments to millions, 2026. URL <https://arxiv.org/abs/2602.02361>.
- [7] NVIDIA Research. CVDP: Comprehensive Verilog design problems benchmark. <https://github.com/NVlabs/CVDP-benchmark>, 2024.
- [8] David Andrews and Saketh Reddy. Spec2GDSBench: Implementation and baseline results for end-to-end chip design evaluation. Technical report, Internal weekly research report, March 2026.
- [9] David Andrews and Saketh Reddy. From individual tasks to full SoC: Redesigning spec2GDSBench around linux boot. Technical report, Internal weekly research report, April 2026.
- [10] David Andrews and Saketh Reddy. Pivoting to AI accelerator benchmarks: An integration ladder for spec2GDSBench. Technical report, Internal weekly research report, April 2026.