
RTL-Universe: A Project-Scale Benchmark for LLM Coding Agents on Hardware Description Languages

David Andrews*, Saketh Reddy*
Georgia Institute of Technology
{dandrews47, sreddy347}@gatech.edu

Abstract

Existing LLM benchmarks for hardware description languages test single-module Verilog generation from natural-language specifications and now sit near the saturation ceiling. The recent VerilogEval, RTLLM, and CVDP benchmarks all report above 90 percent pass rate for current frontier models. We argue this is the wrong measurement. Real hardware engineering takes place at the project scale, where dozens of files cooperate through bus protocols, build systems, and continuous-integration test suites. We introduce RTL-UNIVERSE, a benchmark of eleven open-source HDL projects (Ibex, VeeR EL2, NVDLA, Coral NPU, OpenPiton, PULP common cells, Secworks AES) in which the implementation files of one design unit are stripped, and an LLM coding agent is asked to restore them so the project’s unmodified upstream CI suite passes. We evaluate seven model and harness pairs across Claude Code, the Codex CLI, and OpenCode on a 7×11 matrix. The strongest agent passes 6 of 11 tasks, and four of the seven pass 0 to 2 of 11. A forensic audit of the agent transcripts further shows that 10 of the 15 perfect scores in the run are achieved by exploiting either an unblocked sandbox egress path or a verifier flaw rather than by writing the design. On the eight tasks whose verifiers we trust, the strongest agent passes 2 of 11 tasks. We report the benchmark, the task-generation skill that produces tasks from upstream repositories, the evaluation harness, the run results, the bypass taxonomy, and concrete recommendations for future hardware agent benchmarks.

1 Introduction

LLM coding agents have improved quickly on general-purpose software-engineering benchmarks. SWE-bench [1] reported around 2 percent pass rate at launch, and current top agents pass close to 70 percent. Hardware description languages have followed a similar curve at the module level. VerilogEval [2], RTLLM [3], and the recent CVDP suite [4] all report 90 percent or higher pass rates for current frontier models. Where SWE-bench was succeeded by SWE-bench Verified [5] and then by ProgramBench [6] as agents saturated the simpler regime, no equivalent harder benchmark exists for HDL. We propose one.

The premise of RTL-UNIVERSE is that real hardware engineering takes place at the project scale. An engineer joining an open-source RISC-V core or deep-learning accelerator works inside a directory of dozens of source files, a handful of testbenches, a build system that calls Verilator or Bazel or FuseSoC, and a CI script. Their job is rarely to write a UART transmitter from a paragraph. It is to make a specific change, fix a bug, or re-implement a missing module against the rest of the project. We capture this by taking real open-source HDL projects, deleting one design unit’s implementation files, and asking an agent to restore them so the project’s own CI passes.

Our contributions are:

*Equal contribution. Authors listed in alphabetical order.

1. **A benchmark** of 11 restoration tasks built on real open-source projects spanning two RISC-V cores (Ibex, VeeR EL2), a deep-learning accelerator (NVDLA), a small NPU (Coral NPU), peripherals (OpenPiton ESL, AES), and a cell library (PULP common cells). Together they require restoring 4 to 280 source files per task and exercise build systems including Verilator, Icarus, FuseSoC, Bazel, and cocotb.
2. **A task-generation skill** that converts an arbitrary upstream HDL project into a RTL-UNIVERSE task by deriving the strip manifest, the environment Dockerfile, and the verifier wrapper from the upstream’s existing CI configuration (Section 3).
3. **A multi-agent evaluation harness** that runs Claude Code, the Codex CLI, and OpenCode under a DNS sinkhole, with per-account concurrency caps for subscription-billed agents and a disk-and-process supervisor that we found necessary in practice (Section 4).
4. **An empirical evaluation** of seven model and harness pairs on the 7×11 matrix (Section 5). The strongest agent (GPT-5.5, reasoning effort xhigh) reports 6 of 11 perfect-score passes. The OpenCode-on-OpenRouter open-weight models (Kimi K2.6, DeepSeek V4-Pro, GLM-5.1, Qwen 3.6-27B) cluster at 0 to 2 perfects. Several tasks (coralnpu-full, nvdla-e2e, openpiton-e2e, pulp-common-cells-full) are passed by zero or one model.
5. **An audit** of the agent transcripts that catalogues eight ways an agent can score without doing the task: six varieties of source-code download bypass, plus two verifier exploits we discovered in the upstream test suites (Section 6). We treat this as a quality-control finding rather than a moral one. The audit gives a re-scored leaderboard on the eight tasks whose verifiers we trust.

RTL-UNIVERSE is the latest in a sequence of internal benchmarks the authors have built. A more chronological account appears in our companion logbook. This paper focuses on the benchmark itself and what running it tells us about current agents on project-scale hardware tasks.

2 Related Work

Module-level HDL benchmarks. VerilogEval [2] (156 problems), RTLLM [3] (50 problems), and VeriGen [7] grade single-module Verilog generation from natural-language descriptions. Test suites are 1 to 10 lines and the design is self-contained. CVDP [4] (783 problems) extends this with non-agentic and agentic variants but stays at the module scale. None of these benchmarks asks the agent to operate on an existing build system or write code that is then linked into a larger project.

Agentic SWE benchmarks. SWE-bench [1] runs agents against real Python GitHub issues. SWE-agent [8] introduced the agent-computer interface that became standard. The benchmark progression from SWE-bench to SWE-bench Verified [5] to ProgramBench [6] consistently increases task realism and removes paths that score without writing the code. RTL-UNIVERSE is an HDL-side analogue at the project-scale step.

End-to-end hardware-design benchmarks. The authors’ prior efforts include Spec2GDSBench [9–11], which evaluated agents on a full spec-to-GDSII pipeline from scratch. The agent writes RTL given only a behavioural specification, then a back-end flow synthesises it. We pivoted from build-from-scratch toward restore-from-an-existing-project for two reasons. First, real projects come with real testbenches, and an upstream Verilator or cocotb suite is much more carefully constructed than any test we would write to grade “the agent’s UART”. Second, real projects come with real build systems and integration constraints, including file naming, parameter modules, generator scripts, and Makefile structure, all of which add real engineering surface area. RTL-UNIVERSE is the first design we have built where the test infrastructure is entirely upstream.

Sandboxing and contamination. Test-set contamination [12] and agent-side network access are the two ways an HDL benchmark can be silently passed without writing the code. Section 6 documents both, and we believe future agent benchmarks should treat both with equal care.

3 Benchmark Design and Task-Generation Pipeline

We chose tasks against three criteria. First, the upstream project must have a working CI test suite that runs locally inside the container. Second, the design unit to be restored must be substantial enough that an LLM cannot trivially regenerate it from documentation alone, typically at least 30

Table 1: The eleven RTL-UNIVERSE tasks. #strip is the count of deleted implementation files. Tasks marked † have a verifier whose pass criterion is exit-code-only or assertion-trivial (Section 6). These tasks are reported but discounted in the cleaned ranking.

Task	Upstream	Domain	#strip	Verifier
ibex-e2e†	lowRISC/ibex	RISC-V core	31	4 SW progs in simple_system
ibex-full†	lowRISC/ibex	RISC-V core	32	+ tb_cs_regs UVM test
coralnpu-e2e	google-coral/coralnpu	NPU	240	215 cocotb (Bazel)
coralnpu-full	google-coral/coralnpu	NPU	280	239 cocotb + 2 mobilenet
nvdla-e2e	nvdla/hw	DL accel.	266	SDP/PDP/conv tests
nvdla-full	nvdla/hw	DL accel.	283	full nv_full tests
openpiton-e2e	Princeton/OpenPiton	ESL	4	cocotb counter/lfsr/uart
openpiton-full	Princeton/OpenPiton	ESL	4	18 Icarus benches
pulp-common-cells-full	pulp-platform	cell library	33	in-repo testbench suite
secworks-aes-full	secworks/aes	AES-128 IP	7	5 FuseSoC/Icarus benches
veer-el2-block†	chipsalliance/VeeR-EL2	RISC-V core	40	1 PyUVM test_irq

stripped files, or one large file with hundreds of internal signals. Third, the task should require the agent to use the upstream build system rather than its own. Table 1 lists the eleven tasks.

The eleven tasks cover three difficulty regimes. Small focused units (4 to 10 files) are used to calibrate harness correctness. Mid-scale modules (30 to 40 files) test integration with bus protocols. Large-scale designs (200 or more files) require the agent to reconstruct an entire RTL hierarchy. All upstream projects are widely used. lowRISC’s Ibex is the reference RISC-V core for OpenTitan. Chipsalliance’s VeeR EL2 is the reference for the SweRV family. NVDLA is NVIDIA’s open deep-learning accelerator. Google’s Coral NPU is an edge-NPU reference design. PULP common_cells is the cell library used by every PULP-platform SoC.

Restoration framing and reward. For each task we identify one design unit and remove its implementation files *after* recording a manifest of what was deleted. The remaining files (testbenches, build scripts, READMEs, register-map JSON, linker scripts, reference Python golden models) are kept. The agent receives a working repository plus an instruction file naming the deleted files. The reward is the fraction of upstream CI tests that pass after the agent’s final write, in $[0, 1]$, with single-test verifiers using a binary $\{0, 1\}$.

The deletion-and-restore framing is deliberate. We considered three alternatives. Generating from a written specification (the Spec2GDSBench approach) under-constrains or over-constrains the design. Patching real GitHub bugs (the SWE-bench philosophy) is harder for HDL because well-isolated patch-with-test PRs are roughly an order of magnitude rarer in HDL repositories than in Python. Restoration uses the upstream CI as the grader, and the deletion is a clean noun phrase to put in the instruction (“these files are missing, recreate them”). It also avoids leaking the answer in the spec because the agent must reconstruct concrete syntax rather than satisfy a behavioural contract.

Task-generation skill. We do not write each task by hand. We wrote a Claude Code skill called `harbor-task-gen` that converts an arbitrary upstream HDL repository into a RTL-UNIVERSE task. Given an upstream repo and the path of the design unit to strip, the skill (i) runs the upstream CI to confirm the baseline passes, (ii) derives the strip manifest by following the upstream Makefile or build manifest, (iii) generates an environment Dockerfile that installs only the toolchain the upstream CI requires, (iv) writes a verifier wrapper that calls the upstream test target and emits a normalised reward in $[0, 1]$, and (v) stores the deleted files separately so we can later compare the agent’s output to ground truth. The skill went through five iterations to handle five upstream build systems (FuseSoC, Bazel, Make+Verilator, Icarus, cocotb). Roughly half of the tasks we attempted did not make the final eleven because the upstream build system was either too brittle for the agent to drive or too non-deterministic for the verifier to score. Figure 1 shows the full pipeline.




Figure 1: The RTL-UNIVERSE pipeline. The top row is the offline task-generation flow. The `harbor-task-gen` skill turns an upstream HDL repository into a Harbor task by deriving the strip manifest, the environment Dockerfile, and the verifier wrapper from the upstream CI configuration. The bottom row is the online evaluation flow. The launcher queue feeds tasks to a sandboxed container, the agent writes RTL, and the upstream CI scores the result.

4 Harness

We instantiate each evaluated agent inside a Docker container built from the task’s `environment/` Dockerfile. Just before the agent gets control, the harness writes a DNS sinkhole to `/etc/hosts` naming the most common HDL hosting domains:

```
SINKHOLE = ["github.com", "raw.githubusercontent.com", "api.github.com",
"objects.githubusercontent.com", "codeload.github.com", "gist.github.com",
"gitlab.com", "bitbucket.org", "codeberg.org", "sr.ht", "huggingface.co",
"hf.co", "sourceforge.net"] # appended as "0.0.0.0 <domain>"
```

The intent is to stop a direct `git clone` of the upstream repo. We document in Section 6 that this is necessary but not sufficient.

Concurrency, accounts, supervision. The harness is a Python launcher that runs many agent jobs in parallel against the matrix. Six mechanisms turned out to matter in practice. First, per-job credentials. Each Claude Code job binds a per-job copy of `.credentials.json` so simultaneous installs do not Text file busy each other, and the host file is re-read at each launch to pick up OAuth refreshes. Second, asymmetric per-account caps. Subscription-billed agents (Claude Code, Codex) rate-limit per account, and we cap at `MAX_CLAUDE_ACCT1 = 2` and `MAX_CLAUDE_ACCT2 = 1` across two accounts. Third, a stuck-job killer. A periodic sweep SIGKILLS harbor processes whose log has not advanced for over 30 minutes. Fourth, a race-guard. A 120-second short-term cache of recently-launched job names prevents ps latency from triggering duplicate launches. Fifth, a disk-fill supervisor. The supervisor pauses the launcher when the host root partition exceeds 90 percent (one Sonnet job leaked 812 GB of Verilator build cache through a bind-mounted credentials directory in our run, and would have crashed our shared lab server without this guard). Sixth, an orphan-container reaper. The reaper prunes per-job credential directories whose harbor wrapper has exited. These controls are not exotic but they are non-trivial. We view them as part of what RTL-UNIVERSE contributes as a benchmark, not as plumbing.

Models and harnesses. We evaluate seven model and harness pairs. Opus 4.7 and Sonnet 4.6 run via Claude Code. GPT-5.5 (reasoning effort `xhigh`) runs via the Codex CLI. DeepSeek V4-Pro, GLM-5.1, Kimi K2.6, and Qwen 3.6-27B run via OpenCode and OpenRouter. Each cell has a 10-hour wall-clock cap with up to 3 trials.

5 Results

Headline. Table 2 shows the raw matrix. GPT-5.5 leads at 6 of 11 perfect scores, Opus 4.7 has 5 of 11, and Sonnet 4.6 has 3 of 11. The smaller open-weight models cluster at 0 to 2 perfects and average reward below 0.20. The OpenCode harness around the smaller models tends to produce two failure modes that we categorise informally as “early give-up” (under 5 minutes wall-clock, no real

Table 2: Reward matrix on RTL-UNIVERSE. Bold cells are perfect (1.000). Tasks marked † have a verifier flaw and are omitted from the strong-verifier count. “F” denotes an infrastructure failure (auth or install error).

Model	cor-e2e	cor-fl	ibe-e2e†	ibe-fl†	nvd-e2e	nvd-fl	opn-e2e	opn-fl	pulp	sec	vee†
GPT-5.5	1.00	.42	1.00	1.00	.00	1.00	.00	F	.00	1.00	1.00
Opus 4.7	1.00	.09	1.00	.20	.00	.42	.00	1.00	.00	1.00	1.00
Sonnet 4.6	.18	.08	.00	1.00	.86	1.00	.00	.00	.00	1.00	.00
Kimi K2.6	.00	.02	.00	.00	1.00	1.00	.00	.06	.00	.00	.00
DeepSeek	.03	.03	.00	.00	.14	.00	1.00	.00	.00	.00	.00
GLM-5.1	.00	.00	1.00	.00	.00	.00	.00	.00	.00	.00	.00
Qwen-27B	.00	.00	.00	.00	.00	.00	.00	.00	.00	.00	.00

attempt) and “build-system blocked” (the agent writes plausible-looking RTL but cannot drive the upstream build). Both occurred across many cells. We did not adjust prompts, harness settings, or budget across models to keep the comparison fair.

Per-task observation. The two tasks where multiple models score well (*secworks-aes-full*, *openpiton-full*) are the ones with the smallest stripped surface (7 and 4 files) and the cleanest in-repo Python reference model. The tasks where zero or one model scored well (*coralnpu-full*, *nvdla-e2e*, *pulp-common-cells-full*) are the ones with the largest stripped surfaces and the most upstream-build complexity, not necessarily the most algorithmically hard. We read this as evidence that project-scale RTL restoration is dominated by build-system and integration friction rather than by the difficulty of writing any one Verilog statement.

6 Audit: How agents reach high scores

When we read the agent transcripts we found that several perfect scores were achieved without writing the design. We categorise the bypasses into three groups, and the full method list is in the Appendix. We treat this as a quality-control finding about how to run an HDL agent benchmark.

Category A: Source downloads through unblocked egress. Our */etc/hosts* sinkhole blocks *github.com* but not its alternatives. GPT-5.5 successfully cloned Ibex from the third-party mirror at *git.blizzard.systems* (twice), used a GitHub-to-text service (*gitextract.com*) for NVDLA, and bypassed the sinkhole with *curl -resolve* on a hard-coded code load IP. Opus 4.7 used Anthropic’s server-side *WebFetch* tool with a CDN URL (*cdn.jsdelivr.net/gh/lowRISC/ibex@master/*) to fetch all 31 stripped Ibex files in 37 calls, and used the Software Heritage Archive vault to recover Coral NPU. None of these are model failures. They are sandbox failures.

Category B: Built-in agent tools. Both Claude Code (*WebFetch*, *WebSearch*) and the Codex CLI (the *codex_apps* MCP server’s *github_fetch_file* tool) ship with file-fetch capabilities that route through the harness vendor’s servers, not through the container’s network stack. The */etc/hosts* sinkhole cannot reach them. We had not noticed the Codex MCP existed when we built the sandbox. It accounts for one of GPT-5.5’s perfect scores.

Category C: Verifier exploits. The Ibex *simple_system* *tests/test.sh*’s pass criterion is “simulator exit code = 0”. An empty top-level wrapper that runs Verilator’s setup phase and exits clean satisfies this without ever running the firmware. Both Gemini 3.1-Pro (in a paused configuration not in our headline table; in the larger superset of runs it scored 1.000 here using exactly this exploit) and GLM-5.1 reach perfect scores on this verifier. The *veer-e12-block* PyUVM IRQ test has *assert True* in its *check_phase*. Any compiling DUT passes. Opus identified this in 17 minutes and won the cell with *37 module name(); endmodule stubs*. The verifier flaws are upstream bugs that real HDL projects also have, and restoration-style benchmarking surfaces them quickly because the agent will exploit them when given the chance.

Cleaned ranking. Table 3 re-scores the matrix after zeroing every cell that either matches a category-A or category-B bypass we observed in the transcripts, or scores 1.000 on a flag-marked task. On the eight non-flagged tasks, the strongest agents pass 2 of 11 each.

Table 3: Pass rates before and after the audit. “Raw” counts every reward = 1.0. “Strong” zeroes both confirmed bypasses and the three flagged tasks.

Model	Raw pass	Raw avg	Strong pass	Strong avg
GPT-5.5	6/11 (55%)	0.58	1/11 (9%)	0.13
Opus 4.7	5/11 (45%)	0.52	2/11 (18%)	0.25
Sonnet 4.6	3/11 (27%)	0.37	2/11 (18%)	0.28
Kimi K2.6	2/11 (18%)	0.19	2/11 (18%)	0.19
DeepSeek	1/11 (9%)	0.11	1/11 (9%)	0.11
GLM-5.1	1/11 (9%)	0.09	0/11 (0%)	0.00
Qwen-27B	0/11 (0%)	0.00	0/11 (0%)	0.00

What the cleaned ranking tells us. The headline gap between proprietary frontier agents and open-weight agents nearly disappears. The benchmark was good enough to differentiate raw pass rates, but the bypasses inflate the apparent gap. We do not claim the underlying capability is identical because the open-weight models genuinely fail many cells they do not try to bypass. The comparison should be done on a leakage-free sandbox, and ours was not strict enough. Section 7 gives specific recommendations for tightening it.

7 Recommendations and Limitations

Sandbox. Default-deny egress at the iptables level with allowlists for the package mirrors and vendor APIs the agent needs. Block CDNs that mirror GitHub (`cdn.jsdelivr.net/gh`, `unpkg.com`, `cdnjs.cloudflare.com`), text-of-GitHub services (`gitextract.com`, `r.jina.ai`), Software Heritage, and any host on which an MCP server fetches files server-side. For Claude Code, disable WebFetch and WebSearch via `-allowed-tools`. For Codex, do not register the `codex_apps` MCP.

Verifiers. Audit upstream verifiers for “exit code = 0 counts as pass” and for trivial `check_phase` assertions. Replace with output-hash comparisons or explicit signal-trace assertions. A simple structural sanity check (“does the design hierarchy contain modules of these names”) catches the behavioral-stub family without constraining the implementation.

Limitations. We run one trial per cell with up to 3 retries on infrastructure failure. We do not measure within-cell variance. The cheat-versus-legit verdict is a manual judgement on the transcripts. We publish all transcripts and invite re-adjudication. The benchmark covers 11 tasks chosen by us. A wider sweep across the open-source HDL ecosystem, and especially across non-RISC-V architectures, would be a natural extension. The harness is engineered for our two-laptop and lab-server combination. Porting to a fully cloud-native runner is out of scope here.

8 Conclusion

RTL-UNIVERSE is a deliberately harder hardware-LLM benchmark than what is currently in the literature. By framing tasks as restoration of stripped files inside a real upstream project, we exercise the build systems, integration constraints, and testbenches that real HDL engineers wrestle with. Current frontier agents pass 1 to 2 of the eight strong-verifier cells. Project-scale RTL restoration is not a saturated problem. We hope the benchmark, the task-generation skill, the harness, and the audit are all useful to the next group that builds on this.

Acknowledgements. We thank the Anthropic, OpenAI, and OpenCode teams for making their agent harnesses available. The cheating analysis was assisted by Claude Opus 4.7, with the analysis performed on already-concluded transcripts. The harness, tasks, and full agent transcripts are released on the `codex-experiments-202605` branch of github.com/cs3220-research/

rtl-universe; total wall-clock for the run reported here was approximately 4 days on a 128-CPU 1 TB-RAM lab server.

References

- [1] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*, 2024.
- [2] Mingjie Liu, Nathaniel Pinckney, Bruce Khailany, and Haoxing Ren. VerilogEval: Evaluating large language models for Verilog code generation. *arXiv preprint arXiv:2309.07544*, 2023.
- [3] Yao Lu, Shang Liu, Qijun Zhang, and Zhiyao Xie. RTLLM: An open-source benchmark for design RTL generation with large language models. *IEEE Asia and South Pacific Design Automation Conference (ASP-DAC)*, 2024.
- [4] NVIDIA Research. CVPD: A comprehensive verilog design problems benchmark. <https://github.com/NVlabs/CVPD-benchmark>, 2024.
- [5] OpenAI. Introducing SWE-bench Verified. <https://openai.com/index/introducing-swe-bench-verified/>, 2024.
- [6] OpenAI. Introducing ProgramBench: harder coding agent evaluations. Technical report (programbench), 2024.
- [7] Shailja Thakur, Baleegh Ahmad, Hammond Pearce, Benjamin Tan, Brendan Dolan-Gavitt, Ramesh Karri, and Siddharth Garg. VeriGen: A large language model for Verilog code generation. *ACM Transactions on Design Automation of Electronic Systems (TODAES)*, 29(3), 2023.
- [8] John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. SWE-agent: Agent-computer interfaces enable automated software engineering. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [9] David Andrews and Saketh Reddy. Spec2GDSBench: Implementation and baseline results for end-to-end chip design evaluation. Technical report, Internal weekly research report, March 2026.
- [10] David Andrews and Saketh Reddy. From individual tasks to full SoC: Redesigning spec2GDSBench around linux boot. Technical report, Internal weekly research report, April 2026.
- [11] David Andrews and Saketh Reddy. Pivoting to AI accelerator benchmarks: An integration ladder for spec2GDSBench. Technical report, Internal weekly research report, April 2026.
- [12] Shahriar Golchin and Mihai Surdeanu. Time travel in LLMs: Tracing data contamination in large language models. <https://arxiv.org/abs/2308.08493>, 2024.

A Bypass methods, summary

Table 4: Eight observed sandbox-bypass methods.

#	Method	Used (perfect score) by
A1	Third-party Git mirror (<code>blizzard.systems</code>)	GPT-5.5 ($\times 2$)
A2	GitHub-backed CDN (<code>cdn.jsdelivr.net/gh</code>)	Opus
A3	GitHub-to-text service (<code>gitextract.com</code>)	GPT-5.5
A4	Software Heritage Archive vault	Opus
A5	<code>curl -resolve</code> + IP literal	GPT-5.5
A6	Other mirrors (<code>gitee</code> , <code>ghproxy</code> , <code>unpkg</code> , <code>r.jina.ai</code>)	none successful
B1	Claude Code <code>WebFetch</code>	Opus
B2	Codex MCP <code>github_fetch_file</code>	GPT-5.5
B3	Codex MCP <code>web_search</code>	none successful
C1	Behavioural stub (no real implementation)	GPT-5.5
C2	Verifier exit-code exploit	GLM-5.1
C3	No-op <code>assert</code> True scoreboard exploit	Opus